

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import re
import plotly.express as px
import plotly.graph_objs as go
import plotly.figure_factory as ff
from textblob import TextBlob
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
df=pd.read_csv('aerofit_treadmill.csv')
```

✓ *To explore and Visualise the data*

df

	Product	Age	Gender	Education	MaritalStatus	Usage	Fitness	Income	Miles
0	KP281	18	Male	14	Single	3	4	29562	112
1	KP281	19	Male	15	Single	2	3	31836	75
2	KP281	19	Female	14	Partnered	4	3	30699	66
3	KP281	19	Male	12	Single	3	3	32973	85
4	KP281	20	Male	13	Partnered	4	2	35247	47
...	...	...	...	...	...	...	...	...	...
175	KP781	40	Male	21	Single	6	5	83416	200
176	KP781	42	Male	18	Single	5	4	89641	200
177	KP781	45	Male	16	Single	5	5	90886	160
178	KP781	47	Male	18	Partnered	4	5	104581	120
179	KP781	48	Male	18	Partnered	4	5	95508	180

180 rows × 9 columns


df.head()

	Product	Age	Gender	Education	MaritalStatus	Usage	Fitness	Income	Miles
0	KP281	18	Male	14	Single	3	4	29562	112
1	KP281	19	Male	15	Single	2	3	31836	75
2	KP281	19	Female	14	Partnered	4	3	30699	66
3	KP281	19	Male	12	Single	3	3	32973	85
4	KP281	20	Male	13	Partnered	4	2	35247	47

print(df.info())

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180 entries, 0 to 179
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Product         180 non-null   object
1   Age             180 non-null   int64
2   Gender          180 non-null   object
3   Education       180 non-null   int64
4   MaritalStatus   180 non-null   object
5   Usage           180 non-null   int64
6   Fitness         180 non-null   int64
7   Income          180 non-null   int64
8   Miles           180 non-null   int64
dtypes: int64(6), object(3)
memory usage: 12.8+ KB
None
```

```
plt.figure(figsize=(8, 6))
sns.histplot(data=df, x='Age', bins=20, kde=True)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



```
print(df.describe())
```

	Age	Education	Usage	Fitness	Income \
count	180.000000	180.000000	180.000000	180.000000	180.000000
mean	28.788889	15.572222	3.455556	3.311111	53719.577778
std	6.943498	1.617055	1.084797	0.958869	16506.684226
min	18.000000	12.000000	2.000000	1.000000	29562.000000
25%	24.000000	14.000000	3.000000	3.000000	44058.750000
50%	26.000000	16.000000	3.000000	3.000000	50596.500000
75%	33.000000	16.000000	4.000000	4.000000	58668.000000
max	50.000000	21.000000	7.000000	5.000000	104581.000000


	Miles
count	180.000000
mean	103.194444
std	51.863605
min	21.000000
25%	66.000000
50%	94.000000
75%	114.750000
max	360.000000

```
# Check for missing values
missing_values = df.isnull().sum()
print("Missing Values:")
print(missing_values)
```

```
Missing Values:
Product      0
Age          0
Gender       0
Education    0
MaritalStatus 0
Usage        0
Fitness      0
Income       0
Miles        0
dtype: int64
```

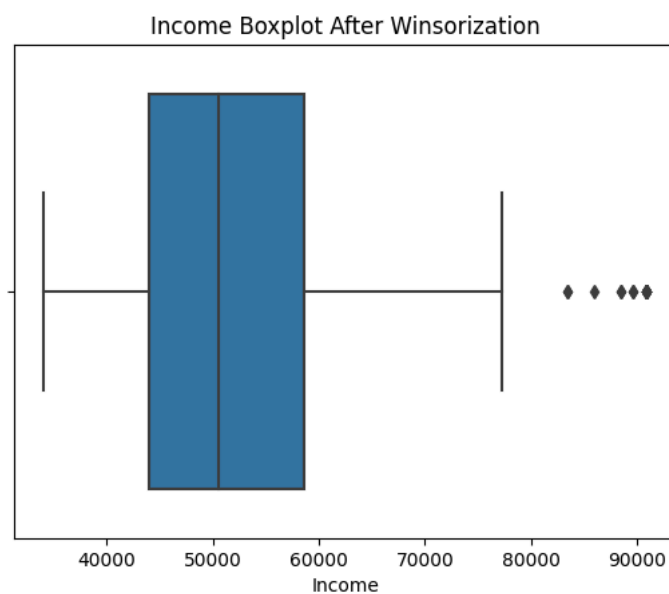
```
# Handling missing values
mean_age = df['Age'].mean()
df['Age'].fillna(mean_age, inplace=True)
```

```
sns.boxplot(x=df['Income'])
plt.title("Income Boxplot")
plt.show()
```



```
from scipy.stats.mstats import winsorize
df['Income'] = winsorize(df['Income'], limits=[0.05, 0.05])
```

```
sns.boxplot(x=df['Income'])
plt.title("Income Boxplot After Winsorization")
plt.show()
```



```
# Summary statistics for numerical attributes
```

```
numerical_attributes = ['Age', 'Usage', 'Income', 'Fitness', 'Miles']
summary_stats = df[numerical_attributes].describe()
print("Summary Statistics for Numerical Attributes:")
print(summary_stats)
```

```
Summary Statistics for Numerical Attributes:
```

	Age	Usage	Income	Fitness	Miles
count	180.000000	180.000000	180.000000	180.000000	180.000000
mean	28.788889	3.455556	53476.800000	3.311111	103.194444
std	6.943498	1.084797	15452.495358	0.958869	51.863605
min	18.000000	2.000000	34110.000000	1.000000	21.000000
25%	24.000000	3.000000	44058.750000	3.000000	66.000000
50%	26.000000	3.000000	50596.500000	3.000000	94.000000
75%	33.000000	4.000000	58668.000000	4.000000	114.750000
max	50.000000	7.000000	90886.000000	5.000000	360.000000

```
# Count of unique values for categorical attributes
```





```
categorical_attributes = ['Gender', 'Education', 'MaritalStatus', 'Product']
unique_counts = df[categorical_attributes].nunique()
print("\nCount of Unique Values for Categorical Attributes:")
print(unique_counts)
```





```
Count of Unique Values for Categorical Attributes:
Gender      2
Education   8
MaritalStatus  2
Product     3
dtype: int64
```


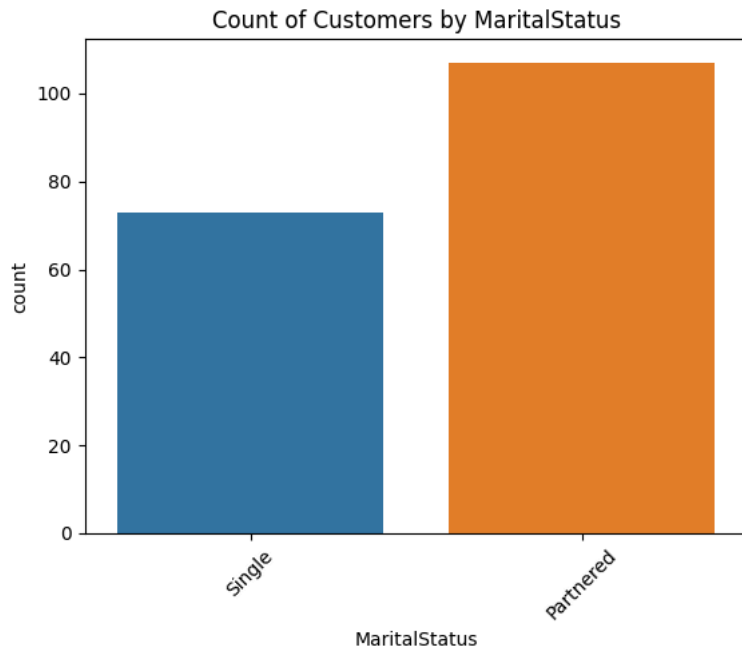
```
# Visualizations to explore attributes
```

```
for attribute in numerical_attributes:
    sns.histplot(data=df, x=attribute, bins=20, kde=True)
    plt.title(f'Distribution of {attribute}')
    plt.show()
```

```
for attribute in categorical_attributes:
    sns.countplot(data=df, x=attribute)
    plt.title(f'Count of Customers by {attribute}')
    plt.xticks(rotation=45)
    plt.show()
```








```
attributes_for_clustering = ['Age', 'Usage', 'Income', 'Fitness', 'Miles']
```

```
scaler = StandardScaler()
data_for_clustering = scaler.fit_transform(df[attributes_for_clustering])
```

```
inertia_values = []
k_values = range(1, 11)
for k in k_values:
    kmeans = KMeans(n_clusters=k, n_init=10, random_state=42)
    kmeans.fit(data_for_clustering)
    inertia_values.append(kmeans.inertia_)
```

```
plt.figure(figsize=(8, 6))
plt.plot(k_values, inertia_values, marker='o', linestyle='-', color='b')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia (Within-Cluster Sum of Squares)')
plt.title('Elbow Method for Optimal K')
plt.grid(True)
plt.show()
```



```

optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, n_init=10, random_state=42)
df['Cluster'] = kmeans.fit_predict(data_for_clustering)
for cluster in range(optimal_k):
    cluster_data = df[df['Cluster'] == cluster]
    plt.scatter(cluster_data['Age'], cluster_data['Income'], label=f'Cluster {cluster + 1}')

plt.xlabel('Age')
plt.ylabel('Income')
plt.title('Customer Clusters')
plt.legend()
plt.grid(True)
plt.show()

```



```

cluster_characteristics = df.groupby('Cluster')[attributes_for_clustering].mean()
print("Cluster Characteristics:")
print(cluster_characteristics)

```

Cluster Characteristics:					
	Age	Usage	Income	Fitness	Miles
Cluster					
0	24.119565	3.141304	44137.804348	2.978261	85.423913
1	30.111111	4.916667	73699.111111	4.777778	183.583333
2	36.134615	3.000000	55999.576923	2.884615	78.980769



```
clustered_data = df.groupby('Cluster')

for cluster, data in clustered_data:
    print(f"Cluster {cluster}:")
    print("-----")
    print("Mean:")
    print(data.mean(numeric_only=True))
    print("\nMedian:")
    print(data.median(numeric_only=True))
    print("\nOther Metrics:")
    print(data.describe())

# to calculate the mean of 'Age' for each cluster
age_means = clustered_data['Age'].mean()
print("Mean Age for Each Cluster:")
print(age_means)
```

```
Cluster 0:
-----
Mean:
Age          24.119565
Education    14.902174
Usage        3.141304
Fitness      2.978261
Income       44137.804348
Miles        85.423913
Cluster      0.000000
dtype: float64


Median:
Age          24.0
Education    14.0
Usage        3.0
Fitness      3.0
Income       45480.0
Miles        85.0
Cluster      0.0
dtype: float64

Other Metrics:
      Age  Education  Usage  Fitness  Income  Miles  \
count  92.000000  92.000000  92.000000  92.000000  92.000000  92.000000
mean   24.119565  14.902174  3.141304  2.978261  44137.804348  85.423913
std     2.676074  1.383258  0.806315  0.695010  7423.448941  27.258725
min     18.000000  12.000000  2.000000  1.000000  34110.000000  38.000000
25%     23.000000  14.000000  3.000000  3.000000  38373.750000  65.500000
50%     24.000000  14.000000  3.000000  3.000000  45480.000000  85.000000
75%     26.000000  16.000000  4.000000  3.000000  49175.250000  106.000000
max     30.000000  21.000000  5.000000  5.000000  69721.000000  170.000000

      Cluster
count      92.0
mean         0.0
std          0.0
min          0.0
25%          0.0
50%          0.0
75%          0.0
max          0.0
Cluster 1:
-----
Mean:
Age          30.111111
Education    17.027778
Usage        4.916667
Fitness      4.777778
Income       73699.111111
Miles       183.583333
Cluster      1.000000
dtype: float64


Median:
Age          28.0
Education    17.0
Usage         5.0
```

```
#bar chart to visualize the distribution of customer segments by age
sns.countplot(data=df, x='Cluster', hue='Age', palette='viridis')
plt.title("Distribution of Customer Segments by Age")
plt.xlabel("Cluster")
plt.ylabel("Count")
plt.show()
```

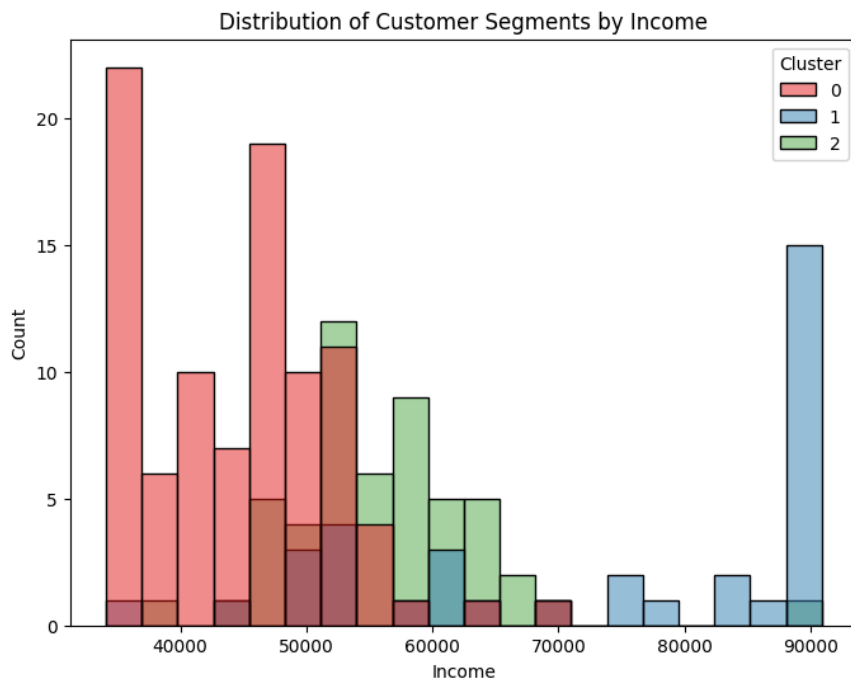


```
# Plot a pie chart to visualize the distribution of customer segments by gender
gender_counts = df['Gender'].value_counts()
plt.pie(gender_counts, labels=gender_counts.index, autopct='%1.1f%%', colors=['skyblue', 'lightcoral'])
plt.title("Distribution of Customer Segments by Gender")
plt.show()
```

Distribution of Customer Segments by Gender



```
# Plot a histogram to visualize the distribution of customer segments by income
plt.figure(figsize=(8, 6))
sns.histplot(data=df, x='Income', hue='Cluster', bins=20, palette='Set1')
plt.title("Distribution of Customer Segments by Income")
plt.xlabel("Income")
plt.ylabel("Count")
plt.show()
```




## Observations

1. There are no missing values in the data.
2. There are 3 unique products in the dataset.
3. KP281 is the most frequent product.
4. Minimum & Maximum age of the person is 18 & 50, mean is 28.79 and 75% of persons have age less than or equal to 33.
5. Most of the people are having 16 years of education i.e., 75% of persons are having education  $\leq 16$  years.
6. Out of 180 data points, 104's gender is Male and rest are the female.
7. Standard deviation for Income & Miles is very high. These variables might have the outliers in it.

## ✓ *checking the structure & characteristics of the dataset*

```
df["Product"].value_counts().plot(kind="bar")
```



```
pd.crosstab(df["Gender"], df["Product"])
```

Product	KP281	KP481	KP781
Gender			
Female	40	29	7
Male	40	31	33

✓ *checking the difference between mean and median*

```


kp281_df = df[df['Product'] == 'KP281']
kp481_df = df[df['Product'] == 'KP481']
kp781_df = df[df['Product'] == 'KP781']
    
```

```

kp281_profile = kp281_df.describe()
kp481_profile = kp481_df.describe()
kp781_profile = kp781_df.describe()
    
```

```

plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='Product', y='Age')
plt.title('Age Distribution by Product')
plt.xlabel('Product')
plt.ylabel('Age')
plt.show()
    
```



```

df['Product'].value_counts()
df['Gender'].value_counts()
df['MaritalStatus'].value_counts()
    
```

```

Partnered    107
Single        73
Name: MaritalStatus, dtype: int64
    
```

```

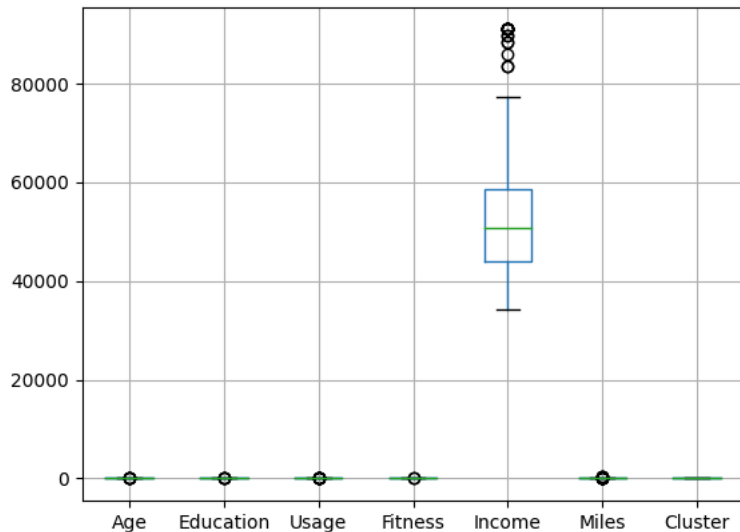
pd.crosstab(df['Product'], df['Gender'])
pd.crosstab(df['Product'], df['MaritalStatus'])
    
```

Product	MaritalStatus	
	Partnered	Single
KP281	48	32
KP481	36	24
KP781	23	17

```

df.boxplot()
    
```

<Axes: >



## ✓ Observations:

From the above boxplots it is quite clear that Age, Education and Usage are having very few outliers while Income and Miles are having more outliers.

```
# Calculate correlation matrix
df.corr(numeric_only=True)
```

	Age	Education	Usage	Fitness	Income	Miles	Cluster
Age	1.000000	0.280496	0.015064	0.061105	0.511992	0.036618	0.751567
Education	0.280496	1.000000	0.395155	0.410581	0.628884	0.307284	0.281891
Usage	0.015064	0.395155	1.000000	0.668606	0.527747	0.759130	0.030955
Fitness	0.061105	0.410581	0.668606	1.000000	0.535939	0.785702	0.056629
Income	0.511992	0.628884	0.527747	0.535939	1.000000	0.554514	0.412124
Miles	0.036618	0.307284	0.759130	0.785702	0.554514	1.000000	0.046590
Cluster	0.751567	0.281891	0.030955	0.056629	0.412124	0.046590	1.000000

```
contingency_kp281 = pd.crosstab(kp281_df['Gender'], kp281_df['MaritalStatus'])
```

```
conditional_prob_kp281 = kp281_df.groupby('Gender')['Fitness'].value_counts(normalize=True)
marginal_prob_gender_kp281 = kp281_df['Gender'].value_counts(normalize=True)
marginal_prob_fitness_kp281 = kp281_df['Fitness'].value_counts(normalize=True)
```

```
conditional_prob_kp281
```

```
Gender  Fitness
Female  3      0.650
        2      0.250
        4      0.075
        5      0.025
Male    3      0.700
        4      0.150
        2      0.100
        1      0.025
        5      0.025
Name: Fitness, dtype: float64
```

```
marginal_prob_gender_kp281
```

```
Male    0.5
Female  0.5
Name: Gender, dtype: float64
```

```
marginal_prob_fitness_kp281
```

```
3    0.6750
2    0.1750
4    0.1125
```

```
5    0.0250
1    0.0125
Name: Fitness, dtype: float64
```

```
conditional_prob_kp481 = kp481_df.groupby('Gender')['Fitness'].value_counts(normalize=True)
marginal_prob_gender_kp481 = kp481_df['Gender'].value_counts(normalize=True)
marginal_prob_fitness_kp481 = kp481_df['Fitness'].value_counts(normalize=True)
```

conditional\_prob\_kp481

```
Gender  Fitness
Female  3      0.620690
        2      0.206897
        4      0.137931
        1      0.034483
Male    3      0.677419
        2      0.193548
        4      0.129032
Name: Fitness, dtype: float64
```

marginal\_prob\_gender\_kp481

```
Male    0.516667
Female  0.483333
Name: Gender, dtype: float64
```

marginal\_prob\_fitness\_kp481

```
3    0.650000
2    0.200000
4    0.133333
1    0.016667
Name: Fitness, dtype: float64
```

```
conditional_prob_kp781 = kp781_df.groupby('Gender')['Fitness'].value_counts(normalize=True)
marginal_prob_gender_kp781 = kp781_df['Gender'].value_counts(normalize=True)
marginal_prob_fitness_kp781 = kp781_df['Fitness'].value_counts(normalize=True)
```

conditional\_prob\_kp781

```
Gender  Fitness
Female  5      0.714286
        3      0.142857
        4      0.142857
Male    5      0.727273
        4      0.181818
        3      0.090909
Name: Fitness, dtype: float64
```

marginal\_prob\_gender\_kp781

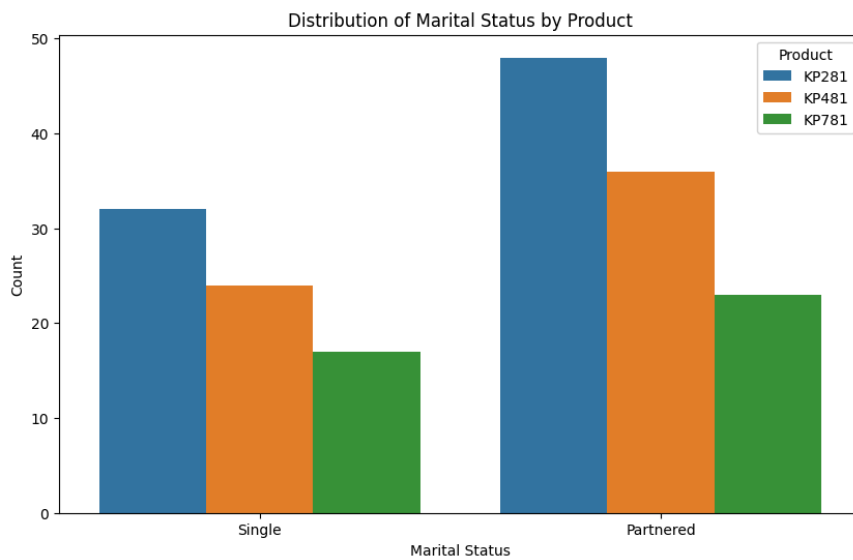
```
Male    0.825
Female  0.175
Name: Gender, dtype: float64
```

marginal\_prob\_fitness\_kp781


```
5    0.725
4    0.175
3    0.100
Name: Fitness, dtype: float64
```

## ✓ ***Countplot to visualize the distribution of marital status by product***

```
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='MaritalStatus', hue='Product')
plt.title('Distribution of Marital Status by Product')
plt.xlabel('Marital Status')
plt.ylabel('Count')
plt.show()
```



```
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='Age', hue='Product', bins=20, kde=True)
plt.title('Age Distribution by Product')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



✓ **Representing the marginal probability like - what percent of customers have purchased KP281, KP481, or KP781 in a table**

```
marginal_prob_table = pd.crosstab(index=df['Product'], columns='Count', normalize='all') * 100
marginal_prob_table.columns = ['Percentage']
marginal_prob_table.reset_index(inplace=True)
print(marginal_prob_table)
```

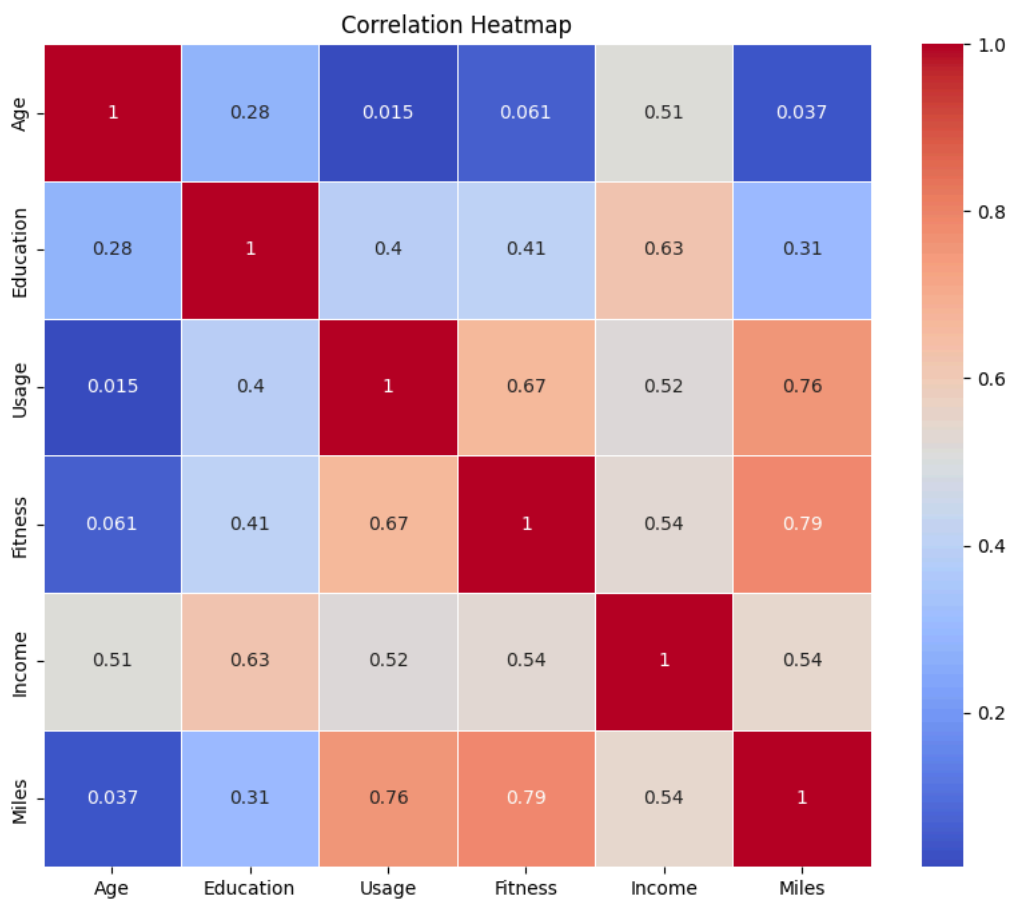
Product	Percentage
0 KP281	44.444444

```
1 KP481 33.333333
2 KP781 22.222222
```

## ✓ Check correlation among different factors using heat maps or pair plots

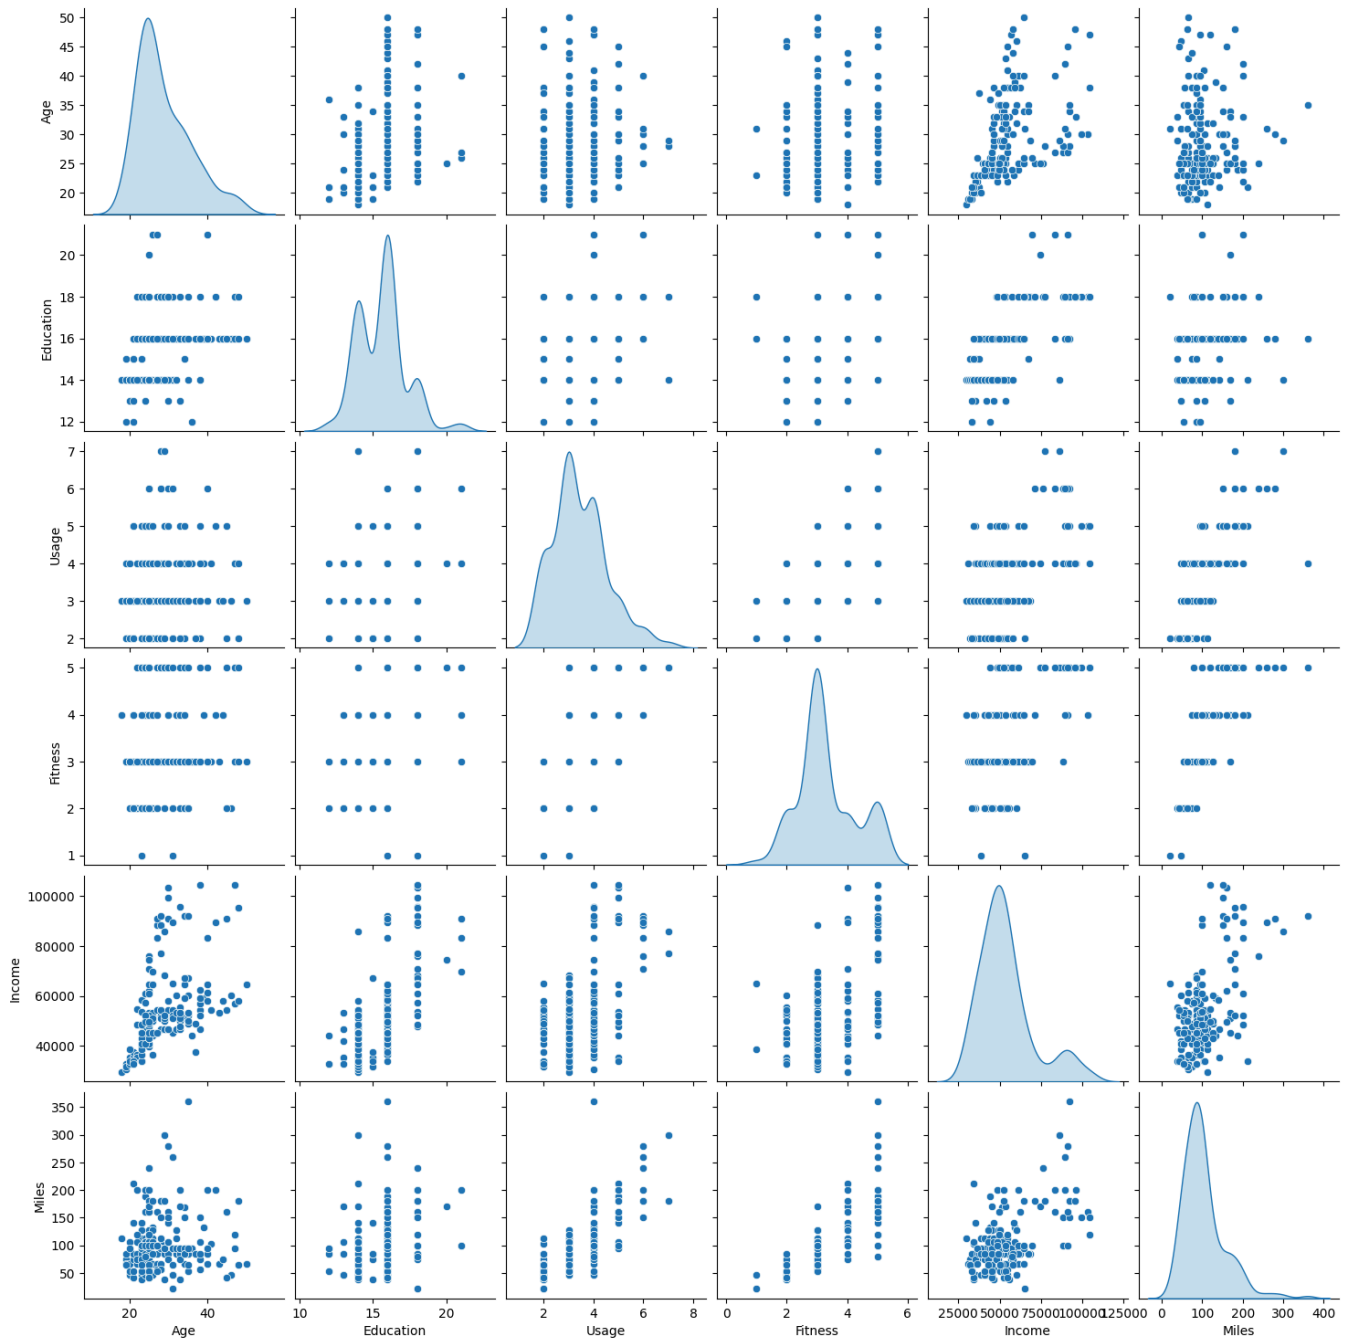
```
correlation_matrix = df.corr(numeric_only=True)
```

```
# Create a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



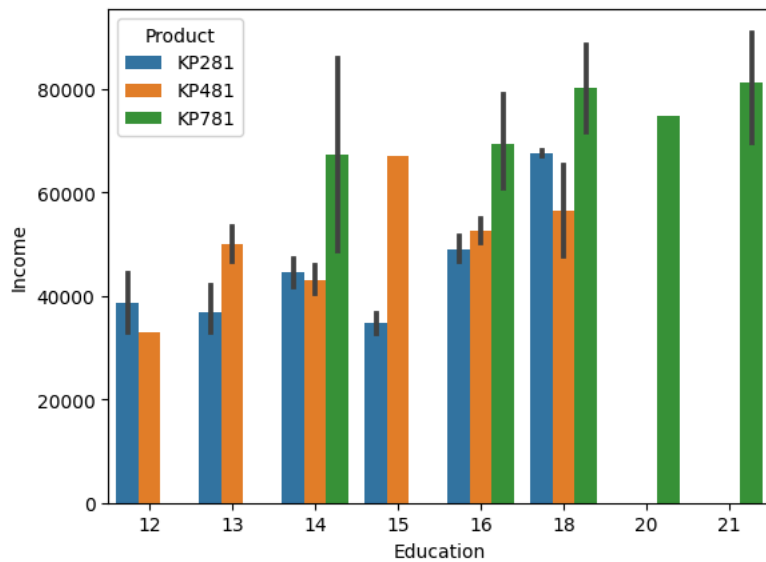
```
# Create a pair plot
sns.pairplot(df, diag_kind='kde')
plt.show()
```





✓ **Recommendation of Treadmill type based on factor Education, Income, Age.**

```
sns.barplot(data = df, x = 'Education', y = 'Income', hue = 'Product')
plt.show()
```



### ✓ *The probability of a male customer buying a KP781 treadmill*

```
# Subset the DataFrame to include only male customers
male_customers = df[df['Gender'] == 'Male']

# Count the number of male customers who purchased KP781
male_kp781_customers = male_customers[male_customers['Product'] == 'KP781'].shape[0]

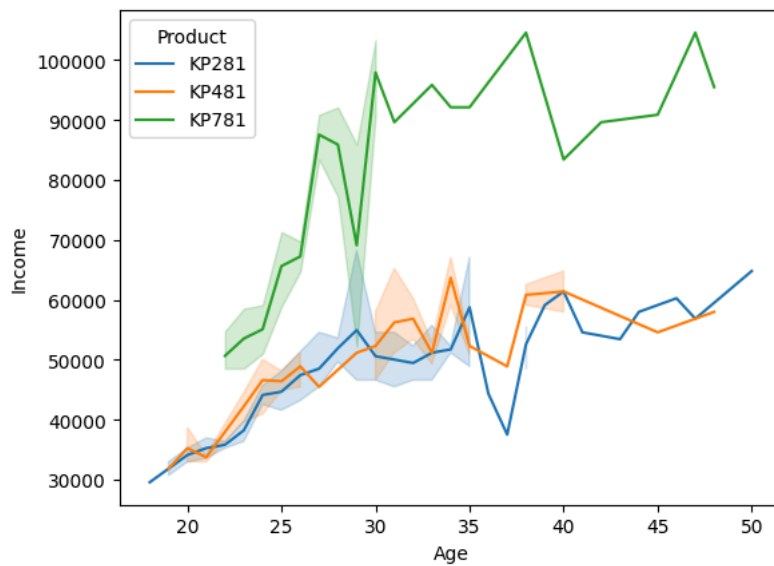
# Calculate the total number of male customers
total_male_customers = male_customers.shape[0]

# Calculate the probability
probability_male_kp781 = male_kp781_customers / total_male_customers

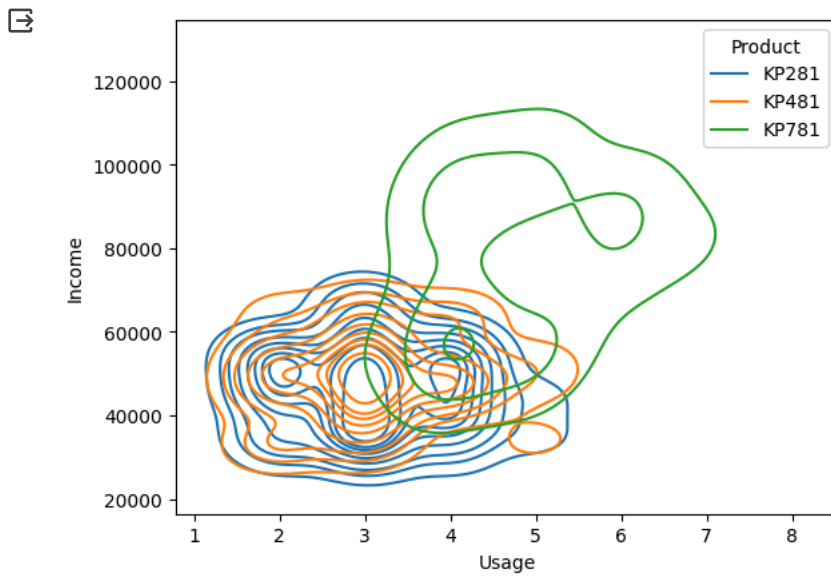

print("Probability of a male customer buying KP781 treadmill:", probability_male_kp781)
```

Probability of a male customer buying KP781 treadmill: 0.3173076923076923

```
sns.lineplot(data = df, x = 'Age', y = 'Income', hue = 'Product')
plt.show()
```



```
sns.boxplot(data = df, x = 'Gender', y = 'Miles', hue = 'Product')
plt.show()
```



```
pd.crosstab(df['Age'], [df['Income'], df['Product']])
```

[illegible]